

# Zadanie domowe: generator obrazów psów

## Autorzy:

- Wojciech Rosenthal, indeks 180458
- Grzegorz Pozorski, indeks 180169

## Informacje:

- Grupa: Uczenie maszynowe 2
- Przedmiot: Uczenie głębokie
- Data: 22.01.2024 r.
- Specjalność: Uczenie maszynowe

## I pobranie danych

```
!curl https://download.microsoft.com/download/3/E/1/3E1C3F21-ECDB-4869-8368-6DEBA77B919F/kagglecatsanddogs_5340.zip -o kagglecatsanddogs_5340.zip
```

% Total		% Received		% Xferd		Average Speed		Time	Time	Time
Current						Dload	Upload	Total	Spent	Left
Speed										
0	0	0	0	0	0	0	0	--:--:--	--:--:--	
--:--:--		0								
0	0	0	0	0	0	0	0	--:--:--	--:--:--	
--:--:--		0								
1	786M	1	8525k	0	0	7345k	0	0:01:49	0:00:01	
0:01:48	7355k									
2	786M	2	17.7M	0	0	8376k	0	0:01:36	0:00:02	
0:01:34	8385k									
3	786M	3	27.2M	0	0	8837k	0	0:01:31	0:00:03	
0:01:28	8843k									
4	786M	4	36.9M	0	0	9084k	0	0:01:28	0:00:04	
0:01:24	9089k									
5	786M	5	46.5M	0	0	9237k	0	0:01:27	0:00:05	
0:01:22	9541k									
7	786M	7	56.2M	0	0	9341k	0	0:01:26	0:00:06	
0:01:20	9804k									
8	786M	8	65.8M	0	0	9418k	0	0:01:25	0:00:07	
0:01:18	9870k									
9	786M	9	75.2M	0	0	9443k	0	0:01:25	0:00:08	

0:01:17 9825k									
10	786M	10	84.9M	0	0	9490k	0	0:01:24	0:00:09
0:01:15 9828k									
12	786M	12	94.5M	0	0	9525k	0	0:01:24	0:00:10
0:01:14 9822k									
12	786M	12	102M	0	0	9368k	0	0:01:25	0:00:11
0:01:14 9403k									
14	786M	14	111M	0	0	9406k	0	0:01:25	0:00:12
0:01:13 9390k									
15	786M	15	121M	0	0	9445k	0	0:01:25	0:00:13
0:01:12 9451k									
16	786M	16	131M	0	0	9476k	0	0:01:25	0:00:14
0:01:11 9451k									
17	786M	17	140M	0	0	9503k	0	0:01:24	0:00:15
0:01:09 9458k									
19	786M	19	150M	0	0	9526k	0	0:01:24	0:00:16
0:01:08 9876k									
20	786M	20	160M	0	0	9547k	0	0:01:24	0:00:17
0:01:07 9890k									
21	786M	21	169M	0	0	9565k	0	0:01:24	0:00:18
0:01:06 9881k									
22	786M	22	179M	0	0	9578k	0	0:01:24	0:00:19
0:01:05 9868k									
24	786M	24	188M	0	0	9593k	0	0:01:23	0:00:20
0:01:03 9867k									
24	786M	24	196M	0	0	9508k	0	0:01:24	0:00:21
0:01:03 9452k									
26	786M	26	206M	0	0	9524k	0	0:01:24	0:00:22
0:01:02 9446k									
27	786M	27	215M	0	0	9536k	0	0:01:24	0:00:23
0:01:01 9431k									
28	786M	28	225M	0	0	9550k	0	0:01:24	0:00:24
0:01:00 9441k									
29	786M	29	234M	0	0	9563k	0	0:01:24	0:00:25
0:00:59 9444k									
31	786M	31	244M	0	0	9573k	0	0:01:24	0:00:26
0:00:58 9845k									
32	786M	32	254M	0	0	9586k	0	0:01:24	0:00:27
0:00:57 9864k									
33	786M	33	263M	0	0	9597k	0	0:01:23	0:00:28
0:00:55 9875k									
34	786M	34	273M	0	0	9606k	0	0:01:23	0:00:29
0:00:54 9880k									
36	786M	36	283M	0	0	9615k	0	0:01:23	0:00:30
0:00:53 9876k									
37	786M	37	292M	0	0	9623k	0	0:01:23	0:00:31
0:00:52 9887k									
38	786M	38	302M	0	0	9630k	0	0:01:23	0:00:32
0:00:51 9869k									

39	786M	39	312M	0	0	9638k	0	0:01:23	0:00:33
0:00:50 9870k									
40	786M	40	319M	0	0	9586k	0	0:01:24	0:00:34
0:00:50 9467k									
41	786M	41	328M	0	0	9555k	0	0:01:24	0:00:35
0:00:49 9196k									
42	786M	42	337M	0	0	9564k	0	0:01:24	0:00:36
0:00:48 9194k									
44	786M	44	347M	0	0	9573k	0	0:01:24	0:00:37
0:00:47 9203k									
45	786M	45	357M	0	0	9579k	0	0:01:24	0:00:38
0:00:46 9190k									
46	786M	46	366M	0	0	9589k	0	0:01:24	0:00:39
0:00:45 9604k									
47	786M	47	376M	0	0	9593k	0	0:01:23	0:00:40
0:00:43 9852k									
49	786M	49	385M	0	0	9600k	0	0:01:23	0:00:41
0:00:42 9855k									
50	786M	50	395M	0	0	9606k	0	0:01:23	0:00:42
0:00:41 9853k									
51	786M	51	405M	0	0	9612k	0	0:01:23	0:00:43
0:00:40 9867k									
52	786M	52	414M	0	0	9613k	0	0:01:23	0:00:44
0:00:39 9804k									
53	786M	53	422M	0	0	9580k	0	0:01:24	0:00:45
0:00:39 9478k									
54	786M	54	432M	0	0	9586k	0	0:01:24	0:00:46
0:00:38 9470k									
56	786M	56	441M	0	0	9592k	0	0:01:23	0:00:47
0:00:36 9473k									
57	786M	57	451M	0	0	9598k	0	0:01:23	0:00:48
0:00:35 9476k									
58	786M	58	461M	0	0	9604k	0	0:01:23	0:00:49
0:00:34 9525k									
59	786M	59	470M	0	0	9610k	0	0:01:23	0:00:50
0:00:33 9881k									
61	786M	61	480M	0	0	9615k	0	0:01:23	0:00:51
0:00:32 9887k									
62	786M	62	490M	0	0	9620k	0	0:01:23	0:00:52
0:00:31 9883k									
63	786M	63	499M	0	0	9624k	0	0:01:23	0:00:53
0:00:30 9876k									
64	786M	64	509M	0	0	9629k	0	0:01:23	0:00:54
0:00:29 9874k									
65	786M	65	516M	0	0	9595k	0	0:01:23	0:00:55
0:00:28 9439k									
66	786M	66	526M	0	0	9600k	0	0:01:23	0:00:56
0:00:27 9443k									
68	786M	68	536M	0	0	9604k	0	0:01:23	0:00:57

0:00:26	9443k								
69	786M	69	545M	0	0	9609k	0	0:01:23	0:00:58
0:00:25	9441k								
70	786M	70	555M	0	0	9612k	0	0:01:23	0:00:59
0:00:24	9420k								
71	786M	71	564M	0	0	9616k	0	0:01:23	0:01:00
0:00:23	9855k								
73	786M	73	574M	0	0	9620k	0	0:01:23	0:01:01
0:00:22	9852k								
74	786M	74	584M	0	0	9624k	0	0:01:23	0:01:02
0:00:21	9851k								
75	786M	75	593M	0	0	9629k	0	0:01:23	0:01:03
0:00:20	9860k								
76	786M	76	603M	0	0	9630k	0	0:01:23	0:01:04
0:00:19	9844k								
77	786M	77	610M	0	0	9601k	0	0:01:23	0:01:05
0:00:18	9422k								
78	786M	78	620M	0	0	9606k	0	0:01:23	0:01:06
0:00:17	9427k								
80	786M	80	630M	0	0	9610k	0	0:01:23	0:01:07
0:00:16	9425k								
81	786M	81	639M	0	0	9614k	0	0:01:23	0:01:08
0:00:15	9426k								
82	786M	82	649M	0	0	9618k	0	0:01:23	0:01:09
0:00:14	9463k								
83	786M	83	659M	0	0	9620k	0	0:01:23	0:01:10
0:00:13	9869k								
85	786M	85	668M	0	0	9624k	0	0:01:23	0:01:11
0:00:12	9872k								
86	786M	86	678M	0	0	9628k	0	0:01:23	0:01:12
0:00:11	9873k								
87	786M	87	688M	0	0	9631k	0	0:01:23	0:01:13
0:00:10	9872k								
88	786M	88	697M	0	0	9635k	0	0:01:23	0:01:14
0:00:09	9871k								
89	786M	89	707M	0	0	9638k	0	0:01:23	0:01:15
0:00:08	9881k								
91	786M	91	717M	0	0	9640k	0	0:01:23	0:01:16
0:00:07	9867k								
92	786M	92	726M	0	0	9644k	0	0:01:23	0:01:17
0:00:06	9872k								
93	786M	93	736M	0	0	9647k	0	0:01:23	0:01:18
0:00:05	9870k								
94	786M	94	743M	0	0	9615k	0	0:01:23	0:01:19
0:00:04	9324k								
95	786M	95	752M	0	0	9618k	0	0:01:23	0:01:20
0:00:03	9328k								
96	786M	96	762M	0	0	9622k	0	0:01:23	0:01:21
0:00:02	9341k								

```

98  786M   98  772M    0    0  9625k    0  0:01:23  0:01:22
0:00:01 9340k
99  786M   99  781M    0    0  9628k    0  0:01:23  0:01:23
--:--:-- 9327k
100 786M  100  786M    0    0  9629k    0  0:01:23  0:01:23
--:--:-- 9876k

```

```
!unzip kagglecatsanddogs_5340.zip
```

```
'unzip' is not recognized as an internal or external command,
operable program or batch file.
```

## II przetwarzanie wstępne

```

import numpy as np
import tensorflow as tf
import cv2
import time
import torch
from PIL import Image
from IPython.display import display
import matplotlib.pyplot as plt
import tensorflow.keras.backend as K

```

```

import os
import random

```

```
print(tf.__version__)
```

```

WARNING:tensorflow:From C:\Users\grzeg\PycharmProjects\pg-workshop\UG\
vae-ug-homework\.venv\lib\site-packages\keras\src\losses.py:2976: The
name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use
tf.compat.v1.losses.sparse_softmax_cross_entropy instead.
2.15.0

```

```

IMAGE_WIDTH=64
IMAGE_HEIGHT=64
IMAGE_SIZE=(IMAGE_WIDTH, IMAGE_HEIGHT)
IMAGE_CHANNELS=3

```

```

dirpath = 'kagglecatsanddogs_5340/PetImages/Dog/'
filenames = [filename for filename in os.listdir(dirpath) if
filename.endswith(".jpg")]
print(f"Total number of images: {len(filenames)}")

```

```
Total number of images: 12500
```

```
sample = dirpath + random.choice(filenamees)
img = Image.open(sample)
display(img)
```



Wykrycie psów na obrazach i przycięcie ich

```
def make_square(image):
    height, width = image.shape[:2]

    max_side = max(height, width)
    top = bottom = left = right = 0

    if height < max_side:
        diff = max_side - height
        top = diff // 2
        bottom = diff - top
    elif width < max_side:
        diff = max_side - width
        left = diff // 2
        right = diff - left

    mean_color = cv2.mean(image)
    mean_color = tuple(map(int, mean_color))
    padded_image = cv2.copyMakeBorder(image, top, bottom, left, right,
cv2.BORDER_CONSTANT, value=mean_color)

    return padded_image

preprocessed_dirpath = 'kagglecatsanddogs_5340/PetImages/Dog-
preprocessed/'

if not os.path.exists(preprocessed_dirpath):
    os.makedirs(preprocessed_dirpath)
```

```
model = torch.hub.load('ultralytics/yolov5', 'yolov5s',  
force_reload=True, trust_repo=True)
```

Downloading: "https://github.com/ultralytics/yolov5/zipball/master" to  
/home/wojtero/.cache/torch/hub/master.zip

Collecting ultralytics

Downloading ultralytics-8.1.3-py3-none-any.whl.metadata (40 kB)  
40.2/40.2 kB 1.0 MB/s eta

0:00:00

matplotlib>=3.3.0 (from ultralytics)

Downloading matplotlib-3.8.2-cp310-cp310-  
manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl.metadata (5.8 kB)  
Requirement already satisfied: numpy>=1.22.2 in

./venv/lib/python3.10/site-packages (from ultralytics) (1.26.3)  
Requirement already satisfied: opencv-python>=4.6.0 in  
./venv/lib/python3.10/site-packages (from ultralytics) (4.9.0.80)

Requirement already satisfied: pillow>=7.1.2 in  
./venv/lib/python3.10/site-packages (from ultralytics) (10.2.0)  
Requirement already satisfied: pyyaml>=5.3.1 in

./venv/lib/python3.10/site-packages (from ultralytics) (6.0.1)  
Requirement already satisfied: requests>=2.23.0 in  
./venv/lib/python3.10/site-packages (from ultralytics) (2.31.0)

Collecting scipy>=1.4.1 (from ultralytics)

Downloading scipy-1.11.4-cp310-cp310-  
manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl.metadata (60 kB)  
60.4/60.4 kB 1.9 MB/s eta

0:00:00

Requirement already satisfied: torch>=1.8.0 in ./venv/lib/python3.10/site-  
packages (from ultralytics) (2.1.2)

Collecting torchvision>=0.9.0 (from ultralytics)

Downloading torchvision-0.16.2-cp310-cp310-  
manylinux1\_x86\_64.whl.metadata (6.6 kB)

Collecting tqdm>=4.64.0 (from ultralytics)

Downloading tqdm-4.66.1-py3-none-any.whl.metadata (57 kB)  
57.6/57.6 kB 4.0 MB/s eta

0:00:00

Requirement already satisfied: psutil in ./venv/lib/python3.10/site-packages  
(from ultralytics) (5.9.8)

Collecting py-cpuinfo (from ultralytics)

Downloading py\_cpuinfo-9.0.0-py3-none-any.whl (22 kB)

Collecting thop>=0.1.1 (from ultralytics)

Downloading thop-0.1.1.post2209072238-py3-none-any.whl (15 kB)  
Requirement already satisfied: pandas>=1.1.4 in  
./venv/lib/python3.10/site-packages (from ultralytics) (2.2.0)

Collecting seaborn>=0.11.0 (from ultralytics)

Downloading seaborn-0.13.1-py3-none-any.whl.metadata (5.4 kB)  
Collecting contourpy>=1.0.1 (from matplotlib>=3.3.0->ultralytics)

Downloading contourpy-1.2.0-cp310-cp310-  
manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl.metadata (5.8 kB)

```
Collecting cyclopy>=0.10 (from matplotlib>=3.3.0->ultralitics)
  Downloading cyclopy-0.12.1-py3-none-any.whl.metadata (3.8 kB)
Collecting fonttools>=4.22.0 (from matplotlib>=3.3.0->ultralitics)
  Downloading fonttools-4.47.2-cp310-cp310-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (157 kB)
157.6/157.6 kB 2.1 MB/s eta
0:00:0000:0100:01
  matplotlib>=3.3.0->ultralitics)
  Downloading kiwisolver-1.4.5-cp310-cp310-
manylinux_2_12_x86_64.manylinux2010_x86_64.whl.metadata (6.4 kB)
Requirement already satisfied: packaging>=20.0 in
./venv/lib/python3.10/site-packages (from matplotlib>=3.3.0-
>ultralitics) (23.2)
Collecting pyparsing>=2.3.1 (from matplotlib>=3.3.0->ultralitics)
  Downloading pyparsing-3.1.1-py3-none-any.whl.metadata (5.1 kB)
Requirement already satisfied: python-dateutil>=2.7 in
./venv/lib/python3.10/site-packages (from matplotlib>=3.3.0-
>ultralitics) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in
./venv/lib/python3.10/site-packages (from pandas>=1.1.4->ultralitics)
(2023.3.post1)
Requirement already satisfied: tzdata>=2022.7 in
./venv/lib/python3.10/site-packages (from pandas>=1.1.4->ultralitics)
(2023.4)
Requirement already satisfied: charset-normalizer<4,>=2 in
./venv/lib/python3.10/site-packages (from requests>=2.23.0-
>ultralitics) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in
./venv/lib/python3.10/site-packages (from requests>=2.23.0-
>ultralitics) (3.6)
Requirement already satisfied: urllib3<3,>=1.21.1 in
./venv/lib/python3.10/site-packages (from requests>=2.23.0-
>ultralitics) (2.1.0)
Requirement already satisfied: certifi>=2017.4.17 in
./venv/lib/python3.10/site-packages (from requests>=2.23.0-
>ultralitics) (2023.11.17)
Requirement already satisfied: filelock in ./venv/lib/python3.10/site-
packages (from torch>=1.8.0->ultralitics) (3.13.1)
Requirement already satisfied: typing-extensions in
./venv/lib/python3.10/site-packages (from torch>=1.8.0->ultralitics)
(4.9.0)
Requirement already satisfied: sympy in ./venv/lib/python3.10/site-
packages (from torch>=1.8.0->ultralitics) (1.12)
Requirement already satisfied: networkx in ./venv/lib/python3.10/site-
packages (from torch>=1.8.0->ultralitics) (3.2.1)
Requirement already satisfied: jinja2 in ./venv/lib/python3.10/site-
packages (from torch>=1.8.0->ultralitics) (3.1.3)
Requirement already satisfied: fsspec in ./venv/lib/python3.10/site-
packages (from torch>=1.8.0->ultralitics) (2023.12.2)
```



Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.1.105 in  
./venv/lib/python3.10/site-packages (from torch>=1.8.0->ultralytics)  
(12.1.105)

Requirement already satisfied: nvidia-cuda-runtime-cu12==12.1.105  
in ./venv/lib/python3.10/site-packages (from torch>=1.8.0-  
>ultralytics) (12.1.105)

Requirement already satisfied: nvidia-cuda-cupti-cu12==12.1.105 in  
./venv/lib/python3.10/site-packages (from torch>=1.8.0->ultralytics)  
(12.1.105)

Requirement already satisfied: nvidia-cudnn-cu12==8.9.2.26 in  
./venv/lib/python3.10/site-packages (from torch>=1.8.0->ultralytics)  
(8.9.2.26)

Requirement already satisfied: nvidia-cublas-cu12==12.1.3.1 in  
./venv/lib/python3.10/site-packages (from torch>=1.8.0->ultralytics)  
(12.1.3.1)

Requirement already satisfied: nvidia-cufft-cu12==11.0.2.54 in  
./venv/lib/python3.10/site-packages (from torch>=1.8.0->ultralytics)  
(11.0.2.54)

Requirement already satisfied: nvidia-curand-cu12==10.3.2.106 in  
./venv/lib/python3.10/site-packages (from torch>=1.8.0->ultralytics)  
(10.3.2.106)

Requirement already satisfied: nvidia-cusolver-cu12==11.4.5.107 in  
./venv/lib/python3.10/site-packages (from torch>=1.8.0->ultralytics)  
(11.4.5.107)

Requirement already satisfied: nvidia-cuspars-cu12==12.1.0.106 in  
./venv/lib/python3.10/site-packages (from torch>=1.8.0->ultralytics)  
(12.1.0.106)

Requirement already satisfied: nvidia-nccl-cu12==2.18.1 in  
./venv/lib/python3.10/site-packages (from torch>=1.8.0->ultralytics)  
(2.18.1)

Requirement already satisfied: nvidia-nvtx-cu12==12.1.105 in  
./venv/lib/python3.10/site-packages (from torch>=1.8.0->ultralytics)  
(12.1.105)

Requirement already satisfied: triton==2.1.0 in  
./venv/lib/python3.10/site-packages (from torch>=1.8.0->ultralytics)  
(2.1.0)

Requirement already satisfied: nvidia-nvjitlink-cu12 in  
./venv/lib/python3.10/site-packages (from nvidia-cusolver-  
cu12==11.4.5.107->torch>=1.8.0->ultralytics) (12.3.101)

Requirement already satisfied: six>=1.5 in ./venv/lib/python3.10/site-  
packages (from python-dateutil>=2.7->matplotlib>=3.3.0->ultralytics)  
(1.16.0)

Requirement already satisfied: MarkupSafe>=2.0 in  
./venv/lib/python3.10/site-packages (from jinja2->torch>=1.8.0-  
>ultralytics) (2.1.4)

Requirement already satisfied: mpmath>=0.19 in  
./venv/lib/python3.10/site-packages (from sympy->torch>=1.8.0-  
>ultralytics) (1.3.0)

Downloading ultralytics-8.1.3-py3-none-any.whl (701 kB)

```

701.1/701.1 kB 7.0 MB/s eta
0:00:00a 0:00:01m
atplotlib-3.8.2-cp310-cp310-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl (11.6 MB)
11.6/11.6 MB 7.2 MB/s eta
0:00:0000:0100:01
anylinux_2_17_x86_64.manylinux2014_x86_64.whl (36.4 MB)
36.4/36.4 MB 34.6 MB/s eta
0:00:0000:0100:01
294.8/294.8 kB 40.7 MB/s eta
0:00:00
anylinux1_x86_64.whl (6.8 MB)
6.8/6.8 MB 40.4 MB/s eta
0:00:00a 0:00:01
-4.66.1-py3-none-any.whl (78 kB)
78.3/78.3 kB 27.2 MB/s eta
0:00:00
anylinux_2_17_x86_64.manylinux2014_x86_64.whl (310 kB)
310.7/310.7 kB 28.7 MB/s eta
0:00:00
anylinux_2_17_x86_64.manylinux2014_x86_64.whl (4.6 MB)
4.6/4.6 MB 34.8 MB/s eta
0:00:00a 0:00:01
anylinux_2_12_x86_64.manylinux2010_x86_64.whl (1.6 MB)
1.6/1.6 MB 38.3 MB/s eta
0:00:0000:01
103.1/103.1 kB 24.3 MB/s eta
0:00:00
, scipy, pyparsing, kiwisolver, fonttools, cyclor, contourpy,
matplotlib, seaborn, torchvision, thop, ultralytics
Successfully installed contourpy-1.2.0 cyclor-0.12.1 fonttools-4.47.2
kiwisolver-1.4.5 matplotlib-3.8.2 py-cpuinfo-9.0.0 pyparsing-3.1.1
scipy-1.11.4 seaborn-0.13.1 thop-0.1.1.post2209072238 torchvision-
0.16.2 tqdm-4.66.1 ultralytics-8.1.3
requirements: Ultralytics requirement ['gitpython>=3.1.30'] not found,
attempting AutoUpdate...
Collecting gitpython>=3.1.30
  Downloading GitPython-3.1.41-py3-none-any.whl.metadata (14 kB)
Collecting gitdb<5,>=4.0.1 (from gitpython>=3.1.30)
  Downloading gitdb-4.0.11-py3-none-any.whl.metadata (1.2 kB)
Collecting smmap<6,>=3.0.1 (from gitdb<5,>=4.0.1->gitpython>=3.1.30)
  Downloading smmap-5.0.1-py3-none-any.whl.metadata (4.3 kB)
Downloading GitPython-3.1.41-py3-none-any.whl (196 kB)
196.4/196.4 kB 1.5 MB/s eta
0:00:00 0:00:010m
62.7/62.7 kB 8.0 MB/s eta
0:00:00
map-5.0.1-py3-none-any.whl (24 kB)
Installing collected packages: smmap, gitdb, gitpython

```

Successfully installed gitdb-4.0.11 gitpython-3.1.41 smmap-5.0.1

requirements: AutoUpdate success ✓ 2.0s, installed 1 package:

['gitpython>=3.1.30']

requirements: ⚠ Restart runtime or rerun command for updates to take effect

YOLOv5 🚀 2024-1-20 Python-3.10.12 torch-2.1.2+cu121 CUDA:0 (NVIDIA GeForce RTX 4080, 16057MiB)

Downloading

<https://github.com/ultralytics/yolov5/releases/download/v7.0/yolov5s.pt> to yolov5s.pt...

100.0%

Fusing layers...

YOLOv5s summary: 213 layers, 7225885 parameters, 0 gradients, 16.4 GFLOPs

Adding AutoShape...

```
sample = 'kagglecatsanddogs_5340/PetImages/Dog/' + '10003.jpg'
```

```
image = cv2.imread(sample, cv2.IMREAD_COLOR)
```

```
results = model(image)
```

```
pd_results = results.pandas().xyxy[0]
```

```
pd_results = pd_results[pd_results['name'] == 'dog']
```

```
pd_results = pd_results[pd_results['confidence'] > 0.5]
```

```
top, bottom, left, right = pd_results[['ymax', 'ymin', 'xmin', 'xmax']].values[0]
```

```
left, right, top, bottom = round(left), round(right), round(top), round(bottom)
```

```
image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
```

```
cropped_image = image[bottom:top, left:right]
```

```
img = Image.fromarray(cropped_image)
```

```
display(img)
```



```
def process_image(filename, dirpath, preprocessed_dirpath):
    try:
        image = cv2.imread(dirpath + filename, cv2.IMREAD_COLOR)
        width, height = image.shape[:2]

        results = model(image)

        pd_results = results.pandas().xyxy[0]
        pd_results = pd_results[pd_results['name'] == 'dog']
        pd_results = pd_results[pd_results['confidence'] > 0.5]
        if len(pd_results) == 0:
            return None
        top, bottom, left, right = pd_results[['ymax', 'ymin', 'xmin',
        'xmax']].values[0]
        left, right, top, bottom = round(left), round(right),
        round(top), round(bottom)

        image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
        cropped_image = image[bottom:top, left:right]

        padded_image = make_square(cropped_image)
        padded_resized_image = cv2.resize(padded_image, IMAGE_SIZE)
```

```

        cv2.imwrite(preprocessed_dirpath + filename,
                    padded_resized_image)

        return width, height
    except Exception as e:
        raise Exception(f"Error processing {filename}: {e}")

mean_size = np.zeros(2)
cnt = 0

start_time = time.time()
for filename in filenames:
    try:
        shape = process_image(filename, dirpath, preprocessed_dirpath)
        if shape:
            mean_size += np.array(shape[:2])
            cnt += 1
    except Exception as e:
        print(f"An error occurred: {e}")

end_time = time.time()
print(f"Total time: {end_time - start_time}")
print(f"Mean size: {mean_size / cnt}.\nTotal count: {cnt}")

```

An error occurred: Error processing 7112.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 6059.jpg: 'NoneType' object has no attribute 'shape'

Corrupt JPEG data: 1403 extraneous bytes before marker 0xd9

An error occurred: Error processing 7459.jpg: 'NoneType' object has no attribute 'shape'

Corrupt JPEG data: 254 extraneous bytes before marker 0xd9

An error occurred: Error processing 10401.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 10158.jpg: 'NoneType' object has no attribute 'shape'

Corrupt JPEG data: 226 extraneous bytes before marker 0xd9

An error occurred: Error processing 1866.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 10797.jpg: 'NoneType' object has no attribute 'shape'

Corrupt JPEG data: 65 extraneous bytes before marker 0xd9

An error occurred: Error processing 3136.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 4367.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 10747.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 11849.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 6718.jpg: 'NoneType' object has no attribute 'shape'

Corrupt JPEG data: 399 extraneous bytes before marker 0xd9

An error occurred: Error processing 5736.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 7133.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 9188.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 5604.jpg: 'NoneType' object has no attribute 'shape'

Corrupt JPEG data: 162 extraneous bytes before marker 0xd9

Warning: unknown JFIF revision number 0.00

An error occurred: Error processing 3288.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 6238.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 8730.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 3588.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 11702.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 7969.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 1308.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 7369.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 2688.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 2384.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 11853.jpg: 'NoneType' object has no attribute 'shape'

An error occurred: Error processing 11410.jpg: 'NoneType' object has no attribute 'shape'

Corrupt JPEG data: 2230 extraneous bytes before marker 0xd9

```
An error occurred: Error processing 11675.jpg: 'NoneType' object has no attribute 'shape'
An error occurred: Error processing 2877.jpg: 'NoneType' object has no attribute 'shape'
Total time: 72.31985402107239
Mean size: [      380.97      420.61].
Total count: 6489
```

Total time: 1510.6200618743896

Mean size: [381 420.6]

Total count: 6487

Wczytanie przetworzonego zbioru danych

```
preprocessed_dirpath = 'kagglecatsanddogs_5340/PetImages/Dog-
preprocessed/'
filenames = [filename for filename in
os.listdir(preprocessed_dirpath)]
print(f"Total number of images: {len(filenames)}")

Total number of images: 6487

X = np.zeros((len(filenames), IMAGE_HEIGHT, IMAGE_WIDTH, 1),
dtype=np.uint8)
for index, filename in enumerate(filenames):
    image = cv2.imread(preprocessed_dirpath + filename,
cv2.IMREAD_GRAYSCALE)
    X[index] = image.reshape(IMAGE_HEIGHT, IMAGE_WIDTH, 1)

X = X / 255.0
X_orig = X.copy()

print(X.shape)

(6487, 64, 64, 1)

ratio = 0.8

X = X_orig.copy()

split_index = round(X.shape[0] * ratio)
X_train_val = X[:split_index]
X_test = X[split_index:]

split_index_train_val = round(X_train_val.shape[0] * ratio)
X_train = X_train_val[:split_index_train_val]
X_val = X_train_val[split_index_train_val:]

print("Training set shape:", X_train.shape)
```

```
print("Validation set shape:", X_val.shape)
print("Test set shape:", X_test.shape)
```

```
Training set shape: (4152, 64, 64, 1)
Validation set shape: (1038, 64, 64, 1)
Test set shape: (1297, 64, 64, 1)
```

```
train_size = X_train.shape[0]
val_size = X_val.shape[0]
test_size = X_test.shape[0]
```

## III Głęboki autokoder wariacyjny

```
# zainspirowane wykładem
```

```
# inspired by
```

```
https://github.com/shashankdhar/VAE-MNIST/blob/master/VAE.py
```

```
def latent_reparam_sampling(args):
    z_mean_layer, z_log_var_layer = args
    batch_size = K.shape(z_mean_layer)[0]
    data_size = K.int_shape(z_mean_layer)[1]
    epsilon = K.random_normal(shape=(batch_size, data_size))
    return z_mean_layer + K.exp(0.5 * z_log_var_layer) * epsilon

class DeepVae(tf.keras.Model):
    def __init__(self, latent_dimension):
        super(DeepVae, self).__init__()

        self.latent_dimension = latent_dimension

        self.encoder = tf.keras.Sequential([
            tf.keras.layers.Input(shape=(64, 64, 1)),
            tf.keras.layers.Conv2D(filters=32, kernel_size=3,
strides=(2, 2), activation='relu'),
            tf.keras.layers.Conv2D(filters=64, kernel_size=3,
strides=(2, 2), activation='relu'),
            tf.keras.layers.Flatten(),
            tf.keras.layers.Dense(self.latent_dimension +
self.latent_dimension),
            ], name='encoder')

        self.decoder = tf.keras.Sequential([
            tf.keras.layers.Input(shape=(self.latent_dimension,)),
            tf.keras.layers.Dense(units=16*16*32,
activation=tf.nn.relu),
            tf.keras.layers.Reshape(target_shape=(16, 16, 32)),
            tf.keras.layers.Conv2DTranspose(filters=64, kernel_size=3,
strides=2, padding='same', activation='relu'),
            tf.keras.layers.Conv2DTranspose(filters=32, kernel_size=3,
```



```

strides=2, padding='same', activation='relu'),
        tf.keras.layers.Conv2DTranspose(filters=1, kernel_size=3,
strides=1, padding='same'),
    ], name='decoder')

```

```

    def get_z_mean_and_z_log_var(self, x):
        z_mean, z_logvar = tf.split(self.encoder(x),
num_or_size_splits=2, axis=1)
        return z_mean, z_logvar

```

```

    def save(self, path):
        self.encoder.save_weights(f'{path}_encoder.h5')
        self.decoder.save_weights(f'{path}_decoder.h5')

```

```

    def load(self, path):
        self.encoder.load_weights(f'{path}_encoder.h5')
        self.decoder.load_weights(f'{path}_decoder.h5')

```

```

def compute_loss(model, x, beta=1.0):
    z_mean, z_log_var = model.get_z_mean_and_z_log_var(x)
    x_p = model.decoder(latent_reparam_sampling((z_mean, z_log_var)))
    loss_reconstruction = K.sum(K.binary_crossentropy(x, x_p), axis=-
1)
    loss_KL = 1 + z_log_var - K.square(z_mean) - K.exp(z_log_var)
    loss_KL = -0.5 * K.sum(loss_KL, axis=-1)
    loss = K.mean(loss_reconstruction) + beta * loss_KL
    return loss, loss_reconstruction, loss_KL

```

@tf.function

```

def train_step(model, x, optimizer, beta=1.0):
    with tf.GradientTape() as tape:
        loss, _, _ = compute_loss(model, x, beta)
        gradients = tape.gradient(loss, model.trainable_variables)
        optimizer.apply_gradients(zip(gradients,
model.trainable_variables))
    return loss

```

```

epochs = 120
annealing_cycles = 4
latent_dimension = 1024

```

```

model = DeepVae(latent_dimension)
model.encoder.summary()
model.decoder.summary()

```

WARNING:tensorflow:From C:\Users\grzeg\PycharmProjects\pg-workshop\UG\vae-ug-homework\.venv\lib\site-packages\keras\src\backend.py:873: The name tf.get\_default\_graph is deprecated. Please use tf.compat.v1.get\_default\_graph instead.

Model: "encoder"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 31, 31, 32)	320
conv2d_1 (Conv2D)	(None, 15, 15, 64)	18496
flatten (Flatten)	(None, 14400)	0
dense (Dense)	(None, 2048)	29493248
=====		
Total params: 29512064 (112.58 MB)		
Trainable params: 29512064 (112.58 MB)		
Non-trainable params: 0 (0.00 Byte)		

Model: "decoder"

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 8192)	8396800
reshape (Reshape)	(None, 16, 16, 32)	0
conv2d_transpose (Conv2DTranspose)	(None, 32, 32, 64)	18496
conv2d_transpose_1 (Conv2DTranspose)	(None, 64, 64, 32)	18464
conv2d_transpose_2 (Conv2DTranspose)	(None, 64, 64, 1)	289
=====		
Total params: 8434049 (32.17 MB)		
Trainable params: 8434049 (32.17 MB)		
Non-trainable params: 0 (0.00 Byte)		

```
optimizer = tf.keras.optimizers.Adam(1e-4)
batch_size = 128
```

```
train_dataset =
(tf.data.Dataset.from_tensor_slices(X_train.astype(np.float32)).shuffle(
train_size).batch(batch_size))
val_dataset =
(tf.data.Dataset.from_tensor_slices(X_val.astype(np.float32)).shuffle(
val_size).batch(batch_size))
test_dataset =
(tf.data.Dataset.from_tensor_slices(X_test.astype(np.float32)).shuffle(
test_size).batch(batch_size))
```

```

annealing_cycle_length = epochs // annealing_cycles
best_loss = None

for epoch in range(1, epochs + 1):
    start_time = time.time()
    beta = min(float(int(epoch) % int(annealing_cycle_length)) /
annealing_cycle_length, 1.0)
    train_loss = tf.keras.metrics.Mean()
    val_loss = tf.keras.metrics.Mean()
    for train_x in train_dataset:
        train_loss(train_step(model, train_x, optimizer, beta))
    for val_x in val_dataset:
        val_loss(compute_loss(model, val_x)[0])
    end_time = time.time()

    if best_loss is None or train_loss.result() < best_loss:
        best_loss = train_loss.result()
        model.save('best')

    print(f'Epoch {epoch}: train_loss={train_loss.result()},
val_loss={val_loss.result()}, beta={beta}, time={end_time -
start_time}')

```

```

Epoch 1: train_loss=2.1104588508605957, val_loss=1.4946001768112183,
beta=0.03333333333333333, time=22.639453649520874
Epoch 2: train_loss=1.0628458261489868, val_loss=1.03101646900177,
beta=0.06666666666666667, time=17.791181087493896
Epoch 3: train_loss=0.939510703086853, val_loss=0.9175639748573303,
beta=0.1, time=18.08599328994751
Epoch 4: train_loss=0.8565689325332642, val_loss=0.8327543139457703,
beta=0.13333333333333333, time=18.602849006652832
Epoch 5: train_loss=0.7947717905044556, val_loss=0.7727011442184448,
beta=0.16666666666666666, time=17.969401121139526
Epoch 6: train_loss=0.7531177401542664, val_loss=0.7386229038238525,
beta=0.2, time=17.85089612007141
Epoch 7: train_loss=0.7295197248458862, val_loss=0.7219509482383728,
beta=0.23333333333333334, time=18.16361904144287
Epoch 8: train_loss=0.71897953748703, val_loss=0.715250551700592,
beta=0.26666666666666666, time=18.21601128578186
Epoch 9: train_loss=0.714522123336792, val_loss=0.7125479578971863,
beta=0.3, time=18.423256635665894
Epoch 10: train_loss=0.7127305269241333, val_loss=0.7114357948303223,
beta=0.3333333333333333, time=18.90371537208557
Epoch 11: train_loss=0.7115664482116699, val_loss=0.7105739116668701,
beta=0.36666666666666664, time=19.064379930496216
Epoch 12: train_loss=0.7108173966407776, val_loss=0.7099395394325256,
beta=0.4, time=19.305054903030396
Epoch 13: train_loss=0.7101265788078308, val_loss=0.709210991859436,
beta=0.43333333333333335, time=19.490837574005127
Epoch 14: train_loss=0.7097421884536743, val_loss=0.7091636061668396,

```

beta=0.4666666666666667, time=20.039698362350464  
Epoch 15: train\_loss=0.7092770338058472, val\_loss=0.7080751657485962,  
beta=0.5, time=20.599578857421875  
Epoch 16: train\_loss=0.7087876796722412, val\_loss=0.7077639698982239,  
beta=0.5333333333333333, time=19.763303756713867  
Epoch 17: train\_loss=0.7083626389503479, val\_loss=0.7080010771751404,  
beta=0.5666666666666667, time=20.24574112892151  
Epoch 18: train\_loss=0.7081132531166077, val\_loss=0.7072455883026123,  
beta=0.6, time=20.274344444274902  
Epoch 19: train\_loss=0.707838773727417, val\_loss=0.7070552110671997,  
beta=0.6333333333333333, time=20.3120174407959  
Epoch 20: train\_loss=0.7071681022644043, val\_loss=0.7065160870552063,  
beta=0.6666666666666666, time=19.220548152923584  
Epoch 21: train\_loss=0.7070402503013611, val\_loss=0.7065091729164124,  
beta=0.7, time=18.740514278411865  
Epoch 22: train\_loss=0.7066516876220703, val\_loss=0.7060900926589966,  
beta=0.7333333333333333, time=19.023975610733032  
Epoch 23: train\_loss=0.7063233852386475, val\_loss=0.7052718997001648,  
beta=0.7666666666666667, time=19.50731921195984  
Epoch 24: train\_loss=0.7061111330986023, val\_loss=0.705385148525238,  
beta=0.8, time=18.497273206710815  
Epoch 25: train\_loss=0.7058065533638, val\_loss=0.7049115300178528,  
beta=0.8333333333333334, time=18.53290867805481  
Epoch 26: train\_loss=0.7053521275520325, val\_loss=0.7047088742256165,  
beta=0.8666666666666667, time=18.554425954818726  
Epoch 27: train\_loss=0.7050511240959167, val\_loss=0.7043158411979675,  
beta=0.9, time=18.99888849258423  
Epoch 28: train\_loss=0.7047553658485413, val\_loss=0.7040398120880127,  
beta=0.9333333333333333, time=19.30444836616516  
Epoch 29: train\_loss=0.7045000791549683, val\_loss=0.7037945985794067,  
beta=0.9666666666666667, time=18.870641708374023  
Epoch 30: train\_loss=0.7040393948554993, val\_loss=0.7034760117530823,  
beta=0.0, time=18.9162437915802  
Epoch 31: train\_loss=0.7037041783332825, val\_loss=0.7028787136077881,  
beta=0.03333333333333333, time=17.863133192062378  
Epoch 32: train\_loss=0.7037039995193481, val\_loss=0.7032443284988403,  
beta=0.06666666666666667, time=17.77161192893982  
Epoch 33: train\_loss=0.7031100392341614, val\_loss=0.7024551630020142,  
beta=0.1, time=17.933035373687744  
Epoch 34: train\_loss=0.7029274702072144, val\_loss=0.7022075057029724,  
beta=0.13333333333333333, time=17.860363245010376  
Epoch 35: train\_loss=0.7027343511581421, val\_loss=0.7020270824432373,  
beta=0.16666666666666666, time=18.80611777305603  
Epoch 36: train\_loss=0.702335695266724, val\_loss=0.7014731764793396,  
beta=0.2, time=18.17232632637024  
Epoch 37: train\_loss=0.7021134495735168, val\_loss=0.7015348672866821,  
beta=0.23333333333333334, time=17.778127670288086  
Epoch 38: train\_loss=0.7019042372703552, val\_loss=0.7008528709411621,  
beta=0.26666666666666666, time=18.157698392868042

Epoch 39: train\_loss=0.7014980316162109, val\_loss=0.700861394405365, beta=0.3, time=17.91987633705139  
Epoch 40: train\_loss=0.7013727426528931, val\_loss=0.700497031211853, beta=0.3333333333333333, time=17.94876456260681  
Epoch 41: train\_loss=0.7010940909385681, val\_loss=0.7003654837608337, beta=0.36666666666666664, time=17.823878526687622  
Epoch 42: train\_loss=0.7008051872253418, val\_loss=0.7003440260887146, beta=0.4, time=17.888017177581787  
Epoch 43: train\_loss=0.7006179690361023, val\_loss=0.7000809907913208, beta=0.43333333333333335, time=17.795143365859985  
Epoch 44: train\_loss=0.7004104256629944, val\_loss=0.6997069120407104, beta=0.46666666666666667, time=17.81314992904663  
Epoch 45: train\_loss=0.7001534700393677, val\_loss=0.6993126273155212, beta=0.5, time=17.948390245437622  
Epoch 46: train\_loss=0.6997759342193604, val\_loss=0.6990453600883484, beta=0.5333333333333333, time=17.730741500854492  
Epoch 47: train\_loss=0.6997024416923523, val\_loss=0.6988089084625244, beta=0.56666666666666667, time=17.84877109527588  
Epoch 48: train\_loss=0.6995878219604492, val\_loss=0.6988939642906189, beta=0.6, time=17.782766103744507  
Epoch 49: train\_loss=0.6992529630661011, val\_loss=0.6988014578819275, beta=0.6333333333333333, time=17.78438425064087  
Epoch 50: train\_loss=0.6989504098892212, val\_loss=0.6982407569885254, beta=0.6666666666666666, time=17.71418023109436  
Epoch 51: train\_loss=0.6986980438232422, val\_loss=0.6980570554733276, beta=0.7, time=17.699395179748535  
Epoch 52: train\_loss=0.6985976099967957, val\_loss=0.697930634021759, beta=0.7333333333333333, time=17.69115447998047  
Epoch 53: train\_loss=0.6984084248542786, val\_loss=0.6976491212844849, beta=0.76666666666666667, time=17.614716291427612  
Epoch 54: train\_loss=0.6982346177101135, val\_loss=0.6977491974830627, beta=0.8, time=17.562248468399048  
Epoch 55: train\_loss=0.6980628967285156, val\_loss=0.6971568465232849, beta=0.8333333333333334, time=17.592220783233643  
Epoch 56: train\_loss=0.6977051496505737, val\_loss=0.6970228552818298, beta=0.86666666666666667, time=17.613053560256958  
Epoch 57: train\_loss=0.6977314352989197, val\_loss=0.6971796154975891, beta=0.9, time=17.6390540599823  
Epoch 58: train\_loss=0.69757080078125, val\_loss=0.6964827179908752, beta=0.9333333333333333, time=19.0875027179718  
Epoch 59: train\_loss=0.6971507668495178, val\_loss=0.6966734528541565, beta=0.96666666666666667, time=19.593557834625244  
Epoch 60: train\_loss=0.6971234083175659, val\_loss=0.6963738203048706, beta=0.0, time=18.186825037002563  
Epoch 61: train\_loss=0.696854829788208, val\_loss=0.6990774869918823, beta=0.03333333333333333, time=18.00300931930542  
Epoch 62: train\_loss=0.696959376335144, val\_loss=0.6969718337059021, beta=0.06666666666666667, time=18.593783140182495  
Epoch 63: train\_loss=0.6965886950492859, val\_loss=0.6963696479797363,

beta=0.1, time=18.351609706878662  
Epoch 64: train\_loss=0.6965106129646301, val\_loss=0.6958004832267761,  
beta=0.1333333333333333, time=18.88807201385498  
Epoch 65: train\_loss=0.6963855624198914, val\_loss=0.6955680847167969,  
beta=0.1666666666666666, time=17.06527018547058  
Epoch 66: train\_loss=0.6961843967437744, val\_loss=0.6956273317337036,  
beta=0.2, time=16.467246294021606  
Epoch 67: train\_loss=0.6961728930473328, val\_loss=0.6956288814544678,  
beta=0.2333333333333334, time=16.38812255859375  
Epoch 68: train\_loss=0.6960312724113464, val\_loss=0.6952754855155945,  
beta=0.2666666666666666, time=16.434219360351562  
Epoch 69: train\_loss=0.6957453489303589, val\_loss=0.6951357126235962,  
beta=0.3, time=16.40023112297058  
Epoch 70: train\_loss=0.6957116723060608, val\_loss=0.6949580907821655,  
beta=0.3333333333333333, time=16.42564845085144  
Epoch 71: train\_loss=0.6955066323280334, val\_loss=0.6950899362564087,  
beta=0.3666666666666666, time=16.39795207977295  
Epoch 72: train\_loss=0.6953009366989136, val\_loss=0.6947097182273865,  
beta=0.4, time=16.50849747657776  
Epoch 73: train\_loss=0.6952520608901978, val\_loss=0.6945493221282959,  
beta=0.4333333333333335, time=16.360748052597046  
Epoch 74: train\_loss=0.6952635645866394, val\_loss=0.6945231556892395,  
beta=0.4666666666666667, time=16.32266855239868  
Epoch 75: train\_loss=0.6950370669364929, val\_loss=0.6943386197090149,  
beta=0.5, time=16.476413249969482  
Epoch 76: train\_loss=0.6949123740196228, val\_loss=0.6941936612129211,  
beta=0.5333333333333333, time=16.531242609024048  
Epoch 77: train\_loss=0.694852888584137, val\_loss=0.694151759147644,  
beta=0.5666666666666667, time=16.545607328414917  
Epoch 78: train\_loss=0.6946910619735718, val\_loss=0.6942607760429382,  
beta=0.6, time=16.44561195373535  
Epoch 79: train\_loss=0.6947070360183716, val\_loss=0.6940546631813049,  
beta=0.6333333333333333, time=16.395132064819336  
Epoch 80: train\_loss=0.6944797039031982, val\_loss=0.693905770778656,  
beta=0.6666666666666666, time=16.42884588241577  
Epoch 81: train\_loss=0.6943473219871521, val\_loss=0.6937331557273865,  
beta=0.7, time=16.67229175567627  
Epoch 82: train\_loss=0.6942225694656372, val\_loss=0.6935571432113647,  
beta=0.7333333333333333, time=16.24709725379944  
Epoch 83: train\_loss=0.6941530108451843, val\_loss=0.6936181783676147,  
beta=0.7666666666666667, time=16.321619749069214  
Epoch 84: train\_loss=0.6941514611244202, val\_loss=0.6936118602752686,  
beta=0.8, time=16.39912486076355  
Epoch 85: train\_loss=0.6941747665405273, val\_loss=0.6933795809745789,  
beta=0.8333333333333334, time=16.599560976028442  
Epoch 86: train\_loss=0.6939951777458191, val\_loss=0.6932218074798584,  
beta=0.8666666666666667, time=17.78993582725525  
Epoch 87: train\_loss=0.6938964128494263, val\_loss=0.6932523846626282,  
beta=0.9, time=17.548035383224487

Epoch 88: train\_loss=0.6937994360923767, val\_loss=0.6933560371398926, beta=0.9333333333333333, time=17.601561784744263  
Epoch 89: train\_loss=0.6935822367668152, val\_loss=0.6929031014442444, beta=0.9666666666666667, time=18.136427640914917  
Epoch 90: train\_loss=0.6935515999794006, val\_loss=0.7665694355964661, beta=0.0, time=18.56167435646057  
Epoch 91: train\_loss=0.6936132311820984, val\_loss=0.6929948329925537, beta=0.03333333333333333, time=18.078909397125244  
Epoch 92: train\_loss=0.6934200525283813, val\_loss=0.6929394006729126, beta=0.06666666666666667, time=17.898712396621704  
Epoch 93: train\_loss=0.6933131217956543, val\_loss=0.6929594874382019, beta=0.1, time=17.60557770729065  
Epoch 94: train\_loss=0.6932167410850525, val\_loss=0.6926475167274475, beta=0.13333333333333333, time=17.733420372009277  
Epoch 95: train\_loss=0.6931201815605164, val\_loss=0.692618191242218, beta=0.16666666666666666, time=17.98951506614685  
Epoch 96: train\_loss=0.6931238174438477, val\_loss=0.6924722790718079, beta=0.2, time=18.018215894699097  
Epoch 97: train\_loss=0.6930025219917297, val\_loss=0.6926053762435913, beta=0.23333333333333334, time=17.777899742126465  
Epoch 98: train\_loss=0.6929662823677063, val\_loss=0.6926819086074829, beta=0.26666666666666666, time=17.92274308204651  
Epoch 99: train\_loss=0.6929201483726501, val\_loss=0.6922724843025208, beta=0.3, time=17.800713300704956  
Epoch 100: train\_loss=0.6928223967552185, val\_loss=0.6923108696937561, beta=0.3333333333333333, time=17.651853322982788  
Epoch 101: train\_loss=0.6927914619445801, val\_loss=0.6922712922096252, beta=0.36666666666666664, time=18.745437622070312  
Epoch 102: train\_loss=0.692696213722229, val\_loss=0.692059338092804, beta=0.4, time=18.283118963241577  
Epoch 103: train\_loss=0.6926394701004028, val\_loss=0.692043125629425, beta=0.43333333333333335, time=17.945343494415283  
Epoch 104: train\_loss=0.6925063729286194, val\_loss=0.6921292543411255, beta=0.46666666666666667, time=17.809961318969727  
Epoch 105: train\_loss=0.6925274133682251, val\_loss=0.6919190287590027, beta=0.5, time=17.77335262298584  
Epoch 106: train\_loss=0.6925604939460754, val\_loss=0.6919451951980591, beta=0.5333333333333333, time=17.853723526000977  
Epoch 107: train\_loss=0.6923950910568237, val\_loss=0.6919437050819397, beta=0.5666666666666667, time=18.353167295455933  
Epoch 108: train\_loss=0.6923657059669495, val\_loss=0.6916803121566772, beta=0.6, time=18.16617465019226  
Epoch 109: train\_loss=0.6923002004623413, val\_loss=0.6917513012886047, beta=0.6333333333333333, time=18.153886318206787  
Epoch 110: train\_loss=0.6921660900115967, val\_loss=0.6915661692619324, beta=0.6666666666666666, time=17.887900590896606  
Epoch 111: train\_loss=0.6922193169593811, val\_loss=0.6916059851646423, beta=0.7, time=18.151943922042847  
Epoch 112: train\_loss=0.6920351982116699, val\_loss=0.6916785836219788,

```

beta=0.7333333333333333, time=17.632816791534424
Epoch 113: train_loss=0.6920760273933411, val_loss=0.6913166642189026,
beta=0.7666666666666667, time=18.24550223350525
Epoch 114: train_loss=0.6919795274734497, val_loss=0.6914474964141846,
beta=0.8, time=17.685124397277832
Epoch 115: train_loss=0.6918575167655945, val_loss=0.691393256187439,
beta=0.8333333333333334, time=17.948457717895508
Epoch 116: train_loss=0.6918394565582275, val_loss=0.6914357542991638,
beta=0.8666666666666667, time=18.496946334838867
Epoch 117: train_loss=0.6918550729751587, val_loss=0.6913012266159058,
beta=0.9, time=18.013478755950928
Epoch 118: train_loss=0.6918106079101562, val_loss=0.6912075281143188,
beta=0.9333333333333333, time=18.08651065826416
Epoch 119: train_loss=0.6917456388473511, val_loss=0.691184937953949,
beta=0.9666666666666667, time=18.133899211883545
Epoch 120: train_loss=0.6916614174842834, val_loss=0.6912009119987488,
beta=0.0, time=17.722911834716797

```

## IV Test modelu

```

model.load('best')

test_loss = tf.keras.metrics.Mean()
for test_x in test_dataset:
    test_loss(compute_loss(model, test_x)[0])
print(f'test_loss={test_loss.result()}')

test_loss=0.6904774308204651

import matplotlib.pyplot as plt

def compare_images(image, generated_image):

    fig, ax = plt.subplots(1,2, figsize=(12,6))

    ax[0].imshow(image, cmap='gray')
    ax[0].set_title('Image')
    ax[1].imshow(generated_image, cmap='gray')
    ax[1].set_title('Generated image')

    plt.show()

def generate_image(image):
    encoded = model.encoder(image.reshape((1, 64, 64, 1)))
    args = tf.split(encoded, num_or_size_splits=2, axis=1)
    encoded = latent_reparam_sampling(args)

    return model.decoder.predict(encoded)[0]

```





```

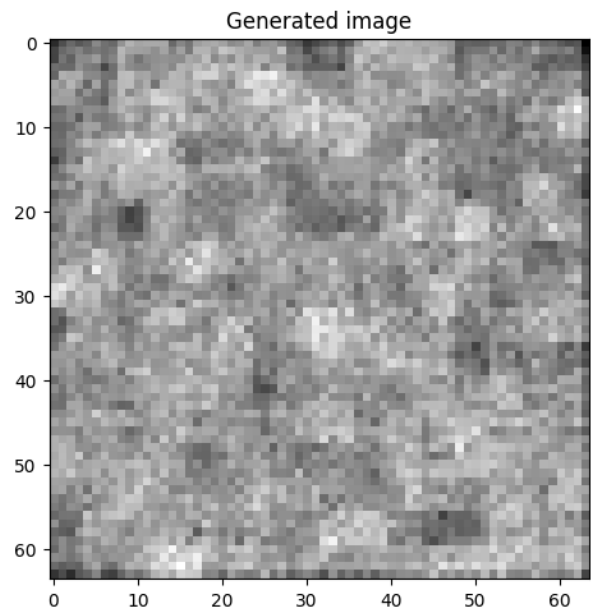

sample_size = 2
sample = np.random.randint(0, X_test.shape[0], sample_size)
for i in range(sample_size):
    decoded = generate_image(X_test[sample[i]])
    denormalized_image = X_test[sample[i]] * 255.0
    compare_images(denormalized_image, decoded)

```

1/1 [=====] - 0s 152ms/step



1/1 [=====] - 0s 18ms/step



## V Generowanie 10 obrazów


```
seed = 33
random_representation = tf.random.uniform((10, 1024), seed=seed)
images = model.decoder(random_representation).numpy().reshape(10, 64,
64, 1)

fig, axes = plt.subplots(nrows=2, ncols=5, figsize=(15, 6))

for i, image in enumerate(images):
    ax = axes[i // 5, i % 5]
    ax.imshow(image.reshape(64, 64, 1), cmap='gray')
    ax.axis('off')

plt.tight_layout()
plt.show()

plt.show()
```



## Additional testing

```
extra_test_dirpath = 'test-dirpath/'
filenames = [filename for filename in os.listdir(extra_test_dirpath)]
print(f"Total number of images: {len(filenames)}")

extra_test_preprocessed_dirpath = 'extra-test-dog-preprocessed/'

if not os.path.exists(extra_test_preprocessed_dirpath):
    os.makedirs(extra_test_preprocessed_dirpath)

start_time = time.time()
for filename in filenames:
    try:
```

```

        process_image(filename, extra_test_dirpath,
extra_test_preprocessed_dirpath)
    except Exception as e:
        print(f"An error occurred: {e}")

end_time = time.time()
print(f"Total time: {end_time - start_time}")

X_test_extra = np.zeros((len(filenames), IMAGE_HEIGHT, IMAGE_WIDTH,
1), dtype=np.uint8)
for index, filename in enumerate(filenames):
    image = cv2.imread(extra_test_preprocessed_dirpath + filename,
cv2.IMREAD_GRAYSCALE)
    X_test_extra[index] = image.reshape(IMAGE_HEIGHT, IMAGE_WIDTH, 1)

X_test_extra = X_test_extra / 255.0

# X_test_extra = X_train[:10]
seed = 33
extra_test_sample_count = 10
images = model.decoder(random_representation).numpy().reshape(10, 64,
64, 1)

fig, axes = plt.subplots(nrows=2, ncols=5, figsize=(15, 6))

for i in range(extra_test_sample_count):
    image = generate_image(X_test_extra[i])
    ax = axes[i // 5, i % 5]
    ax.imshow(image.reshape(64, 64, 1), cmap='gray')
    ax.axis('off')

plt.tight_layout()
plt.show()

plt.show()

1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 19ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 37ms/step
1/1 [=====] - 0s 19ms/step
1/1 [=====] - 0s 17ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 18ms/step

```

